

Do Birds of a Father Flock Together? Examining Religious Homophily Dynamics at a Catholic University

Michael Lee Wood
Brigham Young University

David Hachen
University of Notre Dame

Brandon Sepulvado
NORC at the University of Chicago

Omar Lizardo
UCLA

June 30, 2026

Abstract

Traditional homophily indices often conflate individual preferences with baseline structural opportunities, creating systematic biases when comparing minority and majority groups. We address this by examining religious homophily and friendship choice using longitudinal network data from the NetHealth Cohort Study ($N = 722$ students). Using a dyadic framework with an opportunity offset, we adjust for baseline group sizes and physical and social foci, such as residential proximity and socio-demographic homophily. We find that all religious groups exhibit persistent inbreeding homophily. However, our opportunity-adjusted models reveal that minority religious affiliation (e.g., Protestant and other religions) has a stronger positive effect on active homophily compared to the Catholic majority, supporting Minority Distinctiveness Theory. While cross-sectional baseline models show modest patterns of homophily for secular students, pooled longitudinal analyses indicate that those with 'No Religion' exhibit a statistically significant, durable tendency for same-group affinity over time.

1 Introduction

1.1 Background

How does the numeric composition of a local social environment shape the formation of friendship ties across religious groups? While the reliable tendency for individuals to associate with similar others—homophily—is one of social science’s most well-established empirical generalizations (McPherson et al., 2001), explaining *why* different groups exhibit varying degrees of homophily within a single population remains an open area of research. Although selectively forming ties with similar others can be an individual choice, it is also influenced by structural factors such as propinquity, opportunity structure, and social foci (Rivera et al., 2010; Feld and Grofman, 2009). In this regard, network analysts distinguish between “choice homophily” resulting from individual preferences and “induced homophily” resulting from structural processes such as uneven group sizes or physical proximity (Kossinets and Watts, 2009; McPherson et al., 2001; Marsden, 1988; Hofstra et al., 2017). Furthermore, structural clustering can emerge via “consolidation” or multiform heterogeneity (Blau, 1977a), where preferences for similarity along one dimension (e.g., gender or dormitory) narrow the pool of available alters, indirectly altering apparent homophily along other dimensions like religion.

A subset of structural theories of homophily focuses on how the uneven distribution of group sizes creates motivational (dis)incentives and a biased opportunity structure landscape for forming same-group and out-group ties. In this regard, there are two competing explanations for how relative group size intersects with homophily (and heterophily) generating processes. Under the *Minority Distinctiveness Hypothesis* (MDH), a numeric minority status increases the cognitive salience of that group identity in the context of outgroup numerical dominance, inducing active, deliberate sorting among minority members who band together to preserve their subcultural boundaries (Mehra et al., 1998; Leonard et al., 2008; Marsden, 1988; Sepulvado et al., 2015). Conversely, following Blau’s (1977a) classic treatment, what we might refer to as *Majority Exclusiveness Hypothesis* (MEH) suggests that a dominant majority possesses the structural and social leverage to guard group boundaries, which means that they should exhibit stronger tendencies towards homophily doubly aided by choice (majority members are attractive partners in a context in which they are dominant) and “induced” (there are more members of the majority to choose from) considerations. As Blau observed, “the larger of two groups discriminates more than the smaller against associating with members of the other group,” resulting in higher active homophily among the majority (Blau, 1977a, p. 81).

1.2 Study Setting

An ideal setting in which to examine the MDH and MEH predictions would be a highly specific local context in which a single group holds overwhelming numerical dominance. The University of Notre Dame (ND) presents such a unique social laboratory for the present research. In the *NetHealth Study*’s (Purta et al., 2016) longitudinal cohort, Catholic students represent approximately 73% of the student body, while Protestant (9.9%), non-religious (12.4%), and other religious students (4.6%) constitute small, distinct minorities. However, this setting introduces a profound institutional focus (Feld, 1981; Feld and Grofman, 2009): religion is not merely a demographic attribute; it defines the university’s very identity and physical infrastructure. Crucially, Catholics attending a Catholic university might be more motivated to befriend other Catholics because the environment “summons” their religious identity, increasing its salience (Tavory, 2016). Furthermore, ND features built-in residential and social foci—such as single-sex dormitories, mandatory dorm assignments, and Catholic-oriented campus activities—that could organically channel Catholics into

homophilous friendships without requiring active individual preference for Catholic friends (although such preferences undoubtedly exist in this context).

To evaluate the predictions of the MDH (higher homophily rates among minorities) versus the MEH (higher homophily rates among the majority), it is important to move beyond the limitations of raw homophily rates and aggregate indices. Traditional aggregate measures, such as Krackhardt’s E-I (External-Internal) index (Krackhardt and Stern, 1988), suffer from a well-known artifact: in a bounded network, a numeric minority is structurally forced to have higher outgroup contact, regardless of their personal preferences, a phenomenon also emphasized by Blau (1977b), thus producing a form of “induced heterophily.” Consequently, aggregate comparisons are biased (Everett and Borgatti, 2012). Here, we address that issue by transitioning to a dyadic (tie-level) framework with an opportunity offset. This approach allows us to adjust for baseline group sizes in the data, as well as other important drivers of homophily, such as physical propinquity and social foci, as well as other forms of homophily beyond the focal one in our study.

To this end, this paper leverages 8 waves of longitudinal network data from the *NetHealth Study* to examine the trajectory of religious homophily. We test whether the observed homophily of Protestants and other religious minorities is merely a byproduct of residential sorting and alternative tie-level demographic similarities, or whether it reflects a recurrent, active sorting process driven by identity salience amid Catholic numerical dominance, as predicted by the MDH.

2 Methods and Data

2.1 Data Source and Sample Design

We evaluate the dynamics of religious homophily and friendship choice using longitudinal social network and survey data from the *NetHealth Study* at the University of Notre Dame, a multi-year project funded by the National Institutes of Health. NetHealth tracked a single cohort of undergraduate students from their entry as freshmen in Autumn 2015 through their graduation in Spring 2019. Initial recruitment enlisted around 700 students (yielding a total sample of $N = 722$ participants who took at least one survey), representing roughly 36% of the typical incoming freshman cohort. While this is an exceptionally high coverage rate for an intensive longitudinal ego-network study, the multi-tiered recruitment strategy (which included snowball sampling mechanisms) may introduce some selection bias into homophily estimates if students systematically recruited highly similar alters into the study. The project tracked their social networks, behaviors, and attributes over 8 waves of data collection fielded approximately every 3–4 months.

In each of the eight waves, respondents completed a name-generator survey with the prompt: “We would like to know who you consider to be in your social network. In the spaces below, please list up to 20 people with whom you spend time communicating or interacting.” Beginning with Wave 3, respondents were presented with their previously nominated alters and allowed to add up to 5 of them, yielding a maximum of 25 alters. These name generators, coupled with name interpreters detailing the characteristics of each alter and the nature of the relationship, yielded a unique multiwave, long-term ego-network dataset that constitutes our primary data source. For the purposes of evaluating within-cohort homophily, we restrict alter nominations to those designated as fellow on-campus students.

2.2 Variables

2.2.1 Ego-Level (Participant) Measures

Ego-level attributes are drawn from the NetHealth Basic Survey, which captured demographic and behavioral characteristics across the eight waves. Baseline traits were measured at Wave 1. The primary independent variable is baseline *ego religion*. Respondents were categorized into four groups: Catholic (72.7%, the dominant majority), No Religion (12.3%), Protestant (9.8%), and Other Religion (4.6%). Missing responses account for < 1%. Additional demographic controls include *ego gender*, coded as Woman (51.1%) or Man (47.9%), and *ego race/ethnicity*, categorized into White (65.0%), Latino/a (12.3%), Asian-American (8.7%), Foreign Student (7.1%), African-American (5.7%), and Other (< 1%). To control for individual-level *religious salience (frequency of discussion)* independent of homophily, we use a measure of how frequently the respondent discusses religion (captured in Wave 2). Responses are recorded on a 5-point Likert scale ranging from “Not at all” (4.6%), “Less than 1-2 times a month” (11.8%), “1-2 times a month” (31.9%), “1-2 times a week” (27.4%), to “Three times a week or more” (12.6%).

2.2.2 Alter-Level (Tie) Measures

Alter characteristics and relationship contexts are derived from the NetHealth Network Survey, in which egos provided information about their nominated peers. For the *alter relationship context*, denoting the relation of the alter to the university (e.g., Student, Alum, Prospect), 48.6% of overall nominations were for fellow students. We subset our dyadic data to these alters ($N = 17,452$ student ties). Egos reported the perceived *alter religion* of each nominated fellow student (Catholic: 67.8%; Protestant: 6.8%; No Religion: 7.2%; Other Religion: 2.2%; Missing: 15.9%). This is used to construct our primary outcome variable, a binary indicator of religious homophily (coded 1 if ego and alter share the same religious affiliation, and 0 otherwise). Egos also reported each student *alter gender and race* (e.g., Woman: 52.2%; Man: 47.8%, and White: 78.0%). These allow us to construct binary dyadic indicators for gender homophily and racial/ethnic homophily. To adjust for structurally-induced clustering, egos indicated whether each fellow student shared *physical and social foci (roommates and dorms)*, noting whether the alter was a roommate (13.6% yes) and whether they resided in the same dormitory (34.9% yes). At Notre Dame, the dormitory system constitutes a primary axis of residential and social life. Finally, egos were asked to report how close they felt to each nominated alter on an ordered scale of *friendship closeness (intimacy)*. We construct an ordered numeric indicator for friendship intensity (coded 1 for “Distant”, 2 for “Less Than Close”, 3 for “Merely Close”, and 4 for “Especially Close”). This continuous indicator is added to our multivariate models to control for baseline differences in the homophily of strong versus weak ties.

2.3 Opportunity Offset Formulation

To model the likelihood that a tie is homophilous on religion while adjusting for baseline group proportions, we employ a dyadic logistic regression with an opportunity offset (Scott and Wild, 1991). In highly skewed settings like Notre Dame, where Catholics make up nearly three-quarters of the student body, a traditional regression approach struggles to separate active friendship preferences from simple numeric availability. A Catholic student is highly likely to befriend another Catholic simply because the vast majority of their available peers are Catholic (Blau, 1977b), not necessarily because they are actively seeking out a same-religion friend. The opportunity offset

corrects for this by subtracting the baseline probability of meeting someone of the same religion under conditions of random mixing.

Formally, we specify the model as:

$$\text{logit}(P(\text{Same Religion} = 1)) = \beta_0 + \beta_1 \text{EgoReligion} + \mathbf{X}\beta + \text{opportunity_offset} \quad (1)$$

In this equation, the opportunity_offset is defined as $\log\left(\frac{p_w}{1-p_w}\right)$, where p_w represents the overall proportion of the respondent's religious group in the available pool of peers for that specific wave. By forcing this term into the model with a fixed coefficient of 1.0, we adjust the baseline odds of tie formation. Consequently, the remaining coefficients represent the *excess* tendency to form a homophilous tie above and beyond what the structural opportunity pool dictates.

The model also includes $\mathbf{X}$, a vector of dyadic controls representing alternative mechanisms that might induce homophily without active religious sorting. These include demographic similarities (same gender and same race) as well as physical and social foci (being roommates or living in the same dorm). This specification provides a direct test of the MDH and MEH predictions. A positive β_1 coefficient for a minority group indicates that the group exhibits significantly stronger religious homophily than the Catholic majority, even after accounting for their smaller population size and alternative sorting mechanisms. Such a result would directly support the MDH, demonstrating that numeric minorities actively forge tighter same-group bonds in response to a dominant majority environment.

2.4 Methodological Advantages of the Dyadic Approach

2.4.1 Benefits of the Opportunity Offset

The Krackhardt and Stern (1988) External-Internal (E-I) index is defined as $EI = (E - I)/(E + I)$. Under random mixing, the expected E-I index for group G is a direct linear function of group size p :

$$EI_{\text{expected}} = (1 - p) - p = 1 - 2p \quad (2)$$

Because the baseline shifts with p , raw E-I indices cannot be compared across groups of different sizes to infer preferences. Individual-level aggregations of normalized E-I indices also introduce a severe boundary bias: any minority ego with zero same-religion ties is assigned an EI_{norm} of exactly +1.0. Averaging these individual scores heavily pulls the minority group's average toward +1.0, falsely indicating heterophily. Our dyadic logistic regression with an opportunity offset ($\log(p/(1 - p))$) mitigates these biases. Entering this offset with its coefficient constrained to exactly 1.0 subtracts the baseline opportunity structure. Consequently, the model intercept and group coefficients represent the *excess log-odds of homophilic tie formation beyond what is expected under random mixing*, providing an unbiased, comparative, and control-integrated foundation for social network homophily research.

2.4.2 Why Dyadic Clustered Regression is Preferable to Mixed-Effects Multilevel Models

In social network analysis, egocentric network datasets are inherently hierarchical: nominated friendship ties (Level-1) are nested within individual egos (Level-2 clusters). To correct for this nesting, methodologists frequently recommend mixed-effects logistic regressions:

$$\log \left(\frac{P(Y_{ij} = 1 \mid \gamma_i)}{1 - P(Y_{ij} = 1 \mid \gamma_i)} \right) = \beta_0^C + \beta_1^C \text{EgoReligion}_i + \beta_2^C \mathbf{X}_{ij} + \gamma_i + \text{offset}_{wi} \quad (3)$$

Where $\gamma_i \sim N(0, \sigma^2)$ is a random intercept for ego i , suggesting that multilevel models can be a useful way to model this type of data (Snijders and Baerveldt, 2003). But it can be shown that, for the purposes of answering our questions, the mixed-effects multilevel model is statistically and conceptually less appropriate than a population-average dyadic logistic regression with robust clustered standard errors. First, in a mixed-effects logistic model, the estimated fixed effects (β^C) are conditional: they represent the change in log-odds for a *specific individual* if that individual’s religion were to change. But if religion is a relatively stable baseline trait, as it is in our data, then this is a somewhat less-than-realistic modeling assumption. In contrast, our model with robust clustered standard errors represents a marginal (population-average) model:

$$\log \left(\frac{P(Y_{ij} = 1)}{1 - P(Y_{ij} = 1)} \right) = \beta_0^M + \beta_1^M \text{EgoReligion}_i + \beta_2^M \mathbf{X}_{ij} + \text{offset}_{wi} \quad (4)$$

Substantively, because our core hypothesis concerns group-level differences across the population, the marginal population-average model is the more appropriate approach.

Secondly, the conditional parameters estimated by mixed-effects models are always larger in absolute value than the marginal parameters, indicating that the unobserved ego-level heterogeneity is large and that the random intercepts γ_i act as a “sponge,” absorbing the unobserved cluster-level variance. This collinearity causes the fixed effects in the MLM to shrink toward zero, stripping the model of its statistical power.

Finally, because Protestant and Other Religion students constitute relatively small percentages of the student body, the probability that a minority ego has *exactly zero* same-religion friends is high. When an ego has zero same-religion friends, their Level-1 outcome Y_{ij} has zero variance. In a mixed-effects logistic regression, the maximum likelihood estimate of the random intercept for an ego with zero successes is negative infinity ($\hat{\gamma}_i \rightarrow -\infty$). When the numerical solver attempts to estimate the random effects for these zero-variance clusters, the algorithm fails to find stable starting values and diverges. Our dyadic logistic regression with robust clustered standard errors avoids these numerical and conceptual pitfalls: it is numerically stable, conceptually accurate, and produces statistically defensible standard errors.

3 Results

3.1 Baseline Opportunity Structure Stability

Before analyzing active homophily, it is crucial to establish the stability of the baseline opportunity structure. If the numeric composition of the student body fluctuated wildly over the four years, any changes in tie formation could simply be an artifact of changing demographics rather than evolving social preferences. The wave-specific proportions of available student peers across Waves 3 to 8 are plotted in Figure 1. The extreme flatness of these lines visually confirms that the structural opportunity pool is highly stable, ensuring that the changes in homophily observed in subsequent models reflect genuine behavioral sorting rather than demographic drift.

3.2 Longitudinal Trajectories of Religious Homophily (Yule’s Q)

To address the group-size artifact, we first evaluate the longitudinal trajectory of religious mixing using group-level, baseline-adjusted Yule’s Q across all eight waves of the college cohort (see

Figure 2). Unlike alternative aggregate metrics—such as raw in-group tie percentages, Coleman’s homophily index, or Krackhardt and Stern’s E-I (External-Internal) index—Yule’s Q is independent of relative group sizes (Yule, 1959). Because traditional measures like the E-I index are strictly bounded by the available opportunity pool, they inherently force numeric minorities to exhibit higher rates of heterophily simply due to their small baseline population, distorting cross-group comparisons (for a review mapping segregation measures to theoretical properties, see Bojanowski and Corten, 2014). In contrast, Yule’s Q is a direct transformation of the odds ratio, rendering it margin-free. This ensures that the metric captures active sorting dynamics independent of the severe group size imbalances in our sample, allowing for rigorous and unbiased comparisons between the Catholic majority and various minority groups.¹

Yule’s Q is defined as:

$$Q = \frac{ad - bc}{ad + bc} \quad (5)$$

where a is the number of ingroup ties (e.g., Catholics befriending Catholics), b is the number of outgroup ties from group G (Catholics befriending non-Catholics), c is ties from the outgroup to group G (non-Catholics befriending Catholics), and d is ties strictly between outgroup members (non-Catholics befriending non-Catholics).

The resulting index ranges from -1 to $+1$. A value of 0 indicates perfect random mixing, in which individuals choose friends without regard to religion. A positive value ($Q > 0$) indicates “inbreeding homophily,” meaning individuals actively prefer same-group friends at higher rates than chance would dictate. At the extreme, a Q of $+1$ implies absolute homophily, where group members only ever befriend their own kind. Conversely, a negative value ($Q < 0$) indicates heterophily, meaning individuals actively avoid same-group ties, with -1 indicating that every friendship crosses religious lines.

Our descriptive findings provide strong evidence of *universal inbreeding homophily* ($Q > 0$) for every religious group in every single wave. Over the four years of college, students consistently choose friends who share their religious affiliation at rates that exceed those expected under random mixing.

3.3 Opportunity-Adjusted Cross-Sectional Regressions (Wave 3)

We fit three nested dyadic logistic regression models to Wave 3 student nominations (sophomore-year baseline), adjusting for various tie-level covariates, including physical propinquity (e.g., residing in the same dorm) and gender and race homophily. As Table 1 shows, the nested models reveal that *Protestants and Other Religions exhibit significantly stronger opportunity-adjusted homophily than the Catholic majority*. Crucially, when controlling for roommates, dormitory co-residence, gender homophily, and race homophily matching in Model 2, the Protestant ($\beta = 0.599, p < 0.01$) and Other Religion ($\beta = 1.898, p < 0.001$) homophily coefficients remain highly stable. In contrast, the No Religion coefficient remains positive and statistically significant ($\beta = 0.577, p < 0.05$). Finally, Model 3 introduces a control for individual-level religious salience (frequency of discussing religion). Even after adjusting for religious salience, the active homophily boundaries for Protestant and Other Religion students remain highly significant, indicating that minority distinctiveness is not simply a byproduct of minority students being more religious but reflects a structural group-level boundary.

¹In supplementary analyses, we confirmed that these results are robust to alternative specifications, including point-biserial correlation coefficients, which also show a consistent positive mean tendency for religious homophily after accounting for group-size imbalances.

While such a reduction in nested nonlinear models is often casually interpreted as a mediation effect—suggesting that secular clustering is “explained away” by structural foci like dorms or demographics—this is problematic. In logistic regression, adding covariates rescales the error-term variance, rendering coefficient comparisons across nested models invalid (often called the rescaling or unobserved heterogeneity problem). To determine if physical and demographic foci genuinely mediate the No Religion effect, we conducted a formal causal mediation analysis using the `mediation` package in R (Imai et al., 2010). The formal analysis reveals that the Average Causal Mediation Effects (ACME) for these covariates are all statistically indistinguishable from zero (e.g., $p = 0.32$ for same-dorm co-residence and $p = 0.16$ for race homophily matching). Thus, rather than secular homophily being structurally mediated, students with no religion simply do not exhibit significant active homophily to begin with.

3.4 Pooled Longitudinal Regressions with Robust Clustered Standard Errors

To trace how these associational boundaries evolve longitudinally throughout the college career, we fit a pooled interaction model across Waves 3–8 (Table 2). Because individual students are observed repeatedly, we estimate robust standard errors clustered at the respondent level. The results, summarized in Table 2, show that the No Religion group also exhibits a statistically significant positive effect on active homophily ($\beta = 0.590, p < 0.01$) when pooled across all waves, contrary to the Wave 3 cross-sectional results. Furthermore, we find a counter-intuitive negative effect of same-dorm co-residence on the probability of forming a same-religion tie ($\beta = -0.249, p < 0.05$). We also observe significant racial dynamics: ego-level race effects show that students identifying as African-American form religious homophilous ties at higher rates than White students,² while race homophily (the dyadic tendency to bond with same-race peers) is also a strong, significant predictor of network cohesion ($\beta = 0.433, p < 0.001$), indicating that religious and racial boundaries operate as complementary, rather than competing, axes of social sorting. The results also show that *the heightened homophily between Protestant and Other Religion groups does not change appreciably over time*, as evidenced by highly parallel trends and overlapping confidence intervals across waves.

To evaluate these temporal interaction effects within a non-linear logit framework (Mize, 2019), we calculate and plot the marginal effects for each religious group relative to the Catholic baseline. Because raw probabilities are inherently confounded by unequal group sizes (the group-size artifact), we compute these marginal effects on the probability scale while holding the opportunity offset constant at 0. This analytically simulates an “equal opportunity” scenario in which each group has an identical 50% representation in the population, thereby isolating the excess probability of an active homophilous tie. To further illustrate this stability, Figure 3 plots the wave-specific active homophily coefficients estimated from identical cross-sectional, opportunity-adjusted logistic regressions (similar to Model 1 in Table 1). The resulting plot highlights a clear and durable hierarchy of minority distinctiveness boundaries maintained consistently over the course of college.

4 Discussion and Concluding Remarks

Our results provide strong support for the MDH relative to the MEH predictions. The baseline homophily of the Catholic majority is largely structural—a byproduct of their sheer numeric availability within the university environment, consistent with the more streamlined version of Blau’s

²This finding is likely partially driven by the religious composition of the student body: African-American students in our sample are significantly more likely to be Protestant (29.3%) compared to White students (7.9%), and since Protestant students at this institution exhibit high levels of active homophily, the race coefficient partially captures this religious effect.

(1977b) macro-structural theory. Because Catholics constitute nearly three-quarters of the student body, random mixing naturally yields highly homophilous networks. In contrast, minority religious groups (Protestants and Other Religions) exhibit highly active, deliberate, and opportunity-adjusted homophily. Surrounded by a dominant majority, their minority identity is made highly salient, prompting them to actively build and maintain protective social boundaries to preserve their subcultural beliefs and practices.

Crucially, our revised methodological framework substantially shifts the theoretical conclusions regarding secular students (the "No Religion" group). Previous studies analyzing raw metrics or scaling logistic regressions casually interpreted a drop in the secular homophily coefficient upon the introduction of spatial controls as evidence that secular clustering was merely a byproduct of physical dorm sorting. Instead, the lack of significant active homophily among the No Religion group suggests that secular students lack a sufficiently salient or cohesive shared identity to drive active, boundary-forming friendship preferences. They do not seek each other out. Consistently, the pooled longitudinal model indicates that No Religion students exhibit a statistically significant tendency for active same-group homophily. Additionally, the negative effect of the same-dorm covariate suggests that within the residential dorm units, the 'integrating forces' of university life act to specifically pull students across religious lines, potentially through forced residential mixing that overrides religious preferences.

Furthermore, our longitudinal interaction models, correctly evaluated via marginal effects on the equal-opportunity probability scale, reveal that these active boundary preferences are exceptionally durable throughout the college career. The elevated rates of active homophily among Protestants and Other Religions do not diminish or assimilate over the four years of college. Instead, these groups maintain parallel and remarkably stable trajectories of in-group preference, securing same-religion confidants despite the powerful integrating forces of shared university spaces and randomized residential dorm assignments.

Finally, we note several limitations of the current study. First, there is a substantial selection effect inherent to our sample: non-Catholic students have actively chosen to attend a Catholic university. As a result, they may be less culturally distinct from the Catholic majority than they would be in a secular setting, which may attenuate the boundaries we observe and shift the baseline mechanics of out-group exclusion. Second, because alter attributes are ego-reported, measurement error may be present. In a setting where Catholicism is dominant, respondents who do not know a friend's religion might assume that they are Catholic, potentially artificially inflating Catholic homophily or cross-group ties toward Catholics. Third, we observe a single institutional context with one specific distribution of religious groups. As previous research has shown (Goodreau et al., 2009), the strength of homophily can vary contextually, often exhibiting curvilinear relationships with group proportions rather than strictly linear responses to minority or majority status. Future work extending these findings across multiple institutional contexts could yield greater theoretical specificity regarding the separable effects of group size and group proportion.

References

- Blau, P. M. (1977a). *Inequality and heterogeneity: A primitive theory of social structure*. Free Press.
- Blau, P. M. (1977b). A macrosociological theory of social structure. *American journal of sociology*, 83(1):26–54.
- Bojanowski, M. and Corten, R. (2014). Measuring segregation in social networks. *Social Networks*, 39:14–32.
- Everett, M. G. and Borgatti, S. P. (2012). Categorical attribute based centrality: E-i and g-f centrality. *Social networks*, 34(4):562–569.
- Feld, S. L. (1981). The focused organization of social ties. *American journal of sociology*, 86(5):1015–1035.
- Feld, S. L. and Grofman, B. (2009). Homophily and the focused organization of ties. In Hedstrom, P. and Bearman, P., editors, *Oxford Handbook of Analytical Sociology*, pages 521–543. Oxford University Press.
- Goodreau, S. M., Kitts, J. A., and Morris, M. (2009). Birds of a feather, or friend of a friend? using exponential random graph models to investigate adolescent social networks. *Demography*, 46(1):103–125.
- Hofstra, B., Corten, R., Tubergen, F. v., and Ellison, N. B. (2017). Sources of segregation in social networks: A novel approach using facebook. *American Sociological Review*, 82(3):625–656.
- Imai, K., Keele, L., and Tingley, D. (2010). A general approach to causal mediation analysis. *Psychological methods*, 15(4):309–334.
- Kossinets, G. and Watts, D. J. (2009). Origins of homophily in an evolving social network. *American Journal of Sociology*, 115(2):405–450.
- Krackhardt, D. and Stern, R. N. (1988). Informal networks and organizational crises: An experimental simulation. *Social psychology quarterly*, 51(2):123–140.
- Leonard, A. S., Mehra, A., and Katerberg, R. (2008). The social identity and social networks of ethnic minority groups in organizations: A crucial test of distinctiveness theory. *Journal of Organizational Behavior*, 29(5):573–589.
- Marsden, P. V. (1988). Homogeneity in confiding relations. *Social Networks*, 10(1):57–76.
- McPherson, M., Smith-Lovin, L., and Cook, J. M. (2001). Birds of a feather: Homophily in social networks. *Annual review of sociology*, 27(1):415–444.
- Mehra, A., Kilduff, M., and Brass, D. J. (1998). At the margins: A distinctiveness approach to the social identity and social networks of underrepresented groups. *Academy of management journal*, 41(4):441–452.
- Mize, T. D. (2019). Best practices for estimating, interpreting, and presenting nonlinear interaction effects. *Sociological Science*, 6:81–117.

- Perry, B. L., Pescosolido, B. A., and Borgatti, S. P. (2018). *Egocentric network analysis: Foundations, methods, and models*. Cambridge University Press.
- Purta, R., Mattingly, S., Song, L., Lizardo, O., Hachen, D., Poellabauer, C., and Striegel, A. (2016). Experiences measuring sleep and physical activity patterns across a large college cohort with fitbits. *Proceedings of the 2016 ACM international symposium on wearable computers*, pages 28–35.
- Rivera, M. T., Soderstrom, S. B., and Uzzi, B. (2010). Dynamics of dyads in social networks: Assortative, relational, and proximity mechanisms. *Annual review of sociology*, 36(1):91–115.
- Scott, A. and Wild, C. (1991). Fitting logistic regression models in stratified case-control studies. *Biometrics*, pages 497–510.
- Sepulvado, B., Hachen, D., Penta, M., and Lizardo, O. (2015). Social affiliation from religious disaffiliation: Evidence of selective mixing among youth with no religious preference during the transition to college. *Journal for the Scientific Study of Religion*, 54(4):833–841.
- Snijders, T. A. B. and Baerveldt, C. (2003). A multilevel network study of the effects of delinquent behavior on friendship evolution. *The Journal of mathematical sociology*, 27(2-3):123–151.
- Tavory, I. (2016). *Summoned: Identification and religious life in a Jewish neighborhood*. University of Chicago Press.

Table 1: Nested Dyadic Logistic Regressions Predicting Same-Religion Ties (Wave 3).

	Model 1	Model 2	Model 3
Intercept	0.148* (0.069)	-0.344 (0.320)	0.289 (0.319)
Ego: No Religion	0.698** (0.210)	0.660** (0.220)	0.577* (0.227)
Ego: Other Religion	2.044*** (0.356)	1.957*** (0.388)	1.803*** (0.405)
Ego: Protestant	0.631** (0.193)	0.615** (0.199)	0.603** (0.201)
Ego: Man	—	-0.068 (0.127)	-0.082 (0.129)
Ego: African-American	—	0.312 (0.350)	0.636 (0.373)
Ego: Asian-American	—	-0.233 (0.226)	-0.260 (0.242)
Ego: Foreign Student	—	0.208 (0.364)	-0.163 (0.374)
Ego: Latino/a	—	0.223 (0.218)	0.213 (0.234)
Ego: Other	—	0.926 (0.893)	1.015 (0.892)
Tie: Gender Homophily	—	0.093 (0.185)	0.088 (0.189)
Tie: Race Homophily	—	0.257 (0.174)	0.245 (0.176)
Tie: Roommates	—	-0.171 (0.167)	-0.200 (0.168)
Tie: Same Dorm	—	-0.220 (0.175)	-0.264 (0.178)
Tie: Subjective Closeness	—	-0.034 (0.064)	-0.037 (0.065)
Ego: Religious Salience	—	—	-0.017 (0.046)

Note: * $p < 0.05$; ** $p < 0.01$; *** $p < 0.001$

Table 2: Pooled Longitudinal and Interaction Models (Waves 3-8, Clustered SEs).

	Pooled	Interaction
Intercept	0.043 (0.251)	-0.245* (0.130)
Ego: No Religion	0.576** (0.199)	0.582** (0.201)
Ego: Other Religion	1.778*** (0.531)	1.388*** (0.248)
Ego: Protestant	0.526** (0.184)	0.618** (0.188)
Ego: Man	0.013 (0.055)	0.015 (0.056)
Ego: African-American	0.948*** (0.157)	0.943*** (0.158)
Ego: Asian-American	0.112 (0.106)	0.108 (0.107)
Ego: Foreign Student	0.881*** (0.151)	0.875*** (0.153)
Ego: Latino/a	0.304** (0.100)	0.300** (0.102)
Ego: Other	0.171 (0.560)	0.168 (0.562)
Tie: Gender Homophily	0.115 (0.135)	0.112 (0.079)
Tie: Race Homophily	0.265* (0.118)	0.431*** (0.075)
Tie: Roommates	-0.147 (0.106)	-0.107 (0.075)
Tie: Same Dorm	-0.259* (0.124)	-0.212** (0.076)
Ego: Religious Salience	0.036 (0.061)	0.036 (0.061)
Tie: Subjective Closeness	0.048 (0.029)	0.048 (0.029)
Tie: Wave Number	0.021 (0.019)	0.021 (0.019)
Ego: No Religion x Wave	—	0.076 (0.083)
Ego: Other Religion x Wave	—	0.042 (0.112)
Ego: Protestant x Wave	—	-0.063 (0.076)

Note: * $p < 0.05$; ** $p < 0.01$; *** $p < 0.001$




Figure 1: Religious Opportunity Structure Stability (Proportions of Available Peers, Waves 3-8).




Figure 2: Longitudinal Trajectories of Religious Homophily (Yule's Q) across Waves 3 to 8.




Figure 3: Active (Opportunity-Adjusted) Homophily Coefficients over Time (Waves 3-8).